

Eksploracja danych internetowych

Laboratorium 5

Prowadzący: pracownik UR

Wykonał: Piotr Rojek, pr125159

Zadanie 1

Przy użyciu techniki GridSearchCV wykonaj porównanie modeli klasyfikacyjnych trzech reprezentacji wektorowych:

- Count Vectorizer
- N-Gram
- TF-IDF

Przeprowadź dyskusję wyników.

```
import re
import nltk
import pandas as pd
import string
import warnings

from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

# nltk.download('stopwords')
stopwords = nltk.corpus.stopwords.words("english")
ps = nltk.PorterStemmer()

warnings.filterwarnings( action= "ignore", category=DeprecationWarning)

data = pd.read_csv( filepath_or_buffer: "SMSSpamCollection.tsv", sep='\t', header=None)
data.columns = ['label', 'body_text']

pd.set_option('display.width', 225)
pd.set_option('display.max_columns', 10)
pd.set_option('display.max_colwidth', 50)

def count_punctuation(text): 1 usage
    count = sum([1 for char in text if char in string.punctuation])
    return round(count / (len(text) - text.count(" ")), 3) * 100
```

```
def clean_text(text): 1 usage
    if not isinstance(text, int) and not isinstance(text, float):
        no_punctuation = "".join([char for char in text if char not in string.punctuation])
        lower = no_punctuation.lower()
        tokens = re.split( pattern: r'\W+', lower)
        text = tokens
    text = [word for word in text if word not in stopwords]
    return text
```

```

def steaming(tokenized_text): 1 usage
    text = [ps.stem(word) for word in tokenized_text]
    return text

data['body_text_stemmed'] = data['body_text'].apply(clean_text).apply(steaming)
data['body_text_stemmed'] = (data['body_text_stemmed']
                             .apply(lambda text: ' '.join(text) if isinstance(text, list) else ''))
data['body_len'] = data['body_text'].apply(lambda x: len(x) - x.count(" "))
data['punctuation_%'] = data['body_text'].apply(lambda x: count_punctuation(x))

```

```

# Count Vectorizer
count_vect = CountVectorizer()
x_count = count_vect.fit_transform(data['body_text_stemmed'])
X_count_feat = pd.concat([objs: [data['body_len'], data['punctuation_%']],
                           pd.DataFrame(x_count.toarray(), columns=count_vect.get_feature_names_out())], axis=1)
# print("Count Vectorizer:\n", X_count_feat.head(), "\n")

# N-Gram
ngram_vect = CountVectorizer(ngram_range=(2,2))
x_ngram = ngram_vect.fit_transform(data['body_text_stemmed'])
X_ngram_feat = pd.concat([objs: [data['body_len'], data['punctuation_%']],
                           pd.DataFrame(x_ngram.toarray(), columns=ngram_vect.get_feature_names_out())], axis=1)
# print("N-Gram:\n", X_ngram_feat.head(), "\n")

# TF-IDF
tfidf_vect = TfidfVectorizer()
x_tfidf = tfidf_vect.fit_transform(data['body_text_stemmed'])
X_tfidf_feat = pd.concat([objs: [data['body_len'], data['punctuation_%']],
                           pd.DataFrame(x_tfidf.toarray(), columns=tfidf_vect.get_feature_names_out())], axis=1)
# print("TF-IDF:\n", X_tfidf_feat.head(), "\n")

```

```

# GridSearchCV

rf = RandomForestClassifier()
param = {'n_estimators': [10, 150, 300], 'max_depth': [30, 60, 90, None]}
gs = GridSearchCV(rf, param, cv=5, n_jobs=-1)

gs_fit_count = gs.fit(X_count_feat, data['label'])
print("Count:\n", pd.DataFrame(gs_fit_count.cv_results_).sort_values(by: 'mean_test_score', ascending=False)[0:5])

gs_fit_ngram = gs.fit(X_ngram_feat, data['label'])
print("N-Gram:\n", pd.DataFrame(gs_fit_ngram.cv_results_).sort_values(by: 'mean_test_score', ascending=False)[0:5])

gs_fit_tfidf = gs.fit(X_tfidf_feat, data['label'])
print("TF-IDF:\n", pd.DataFrame(gs_fit_tfidf.cv_results_).sort_values(by: 'mean_test_score', ascending=False)[0:5])

print("Najlepszy mean_test_score:")

gs.fit(X_count_feat, data['label'])
print("Count:", round(gs.best_score_, 5))

gs.fit(X_ngram_feat, data['label'])
print("N-Gram:", round(gs.best_score_, 5))

gs.fit(X_tfidf_feat, data['label'])
print("TF-IDF:", round(gs.best_score_, 5))

```

Testy:

	Count:	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	...	split3_test_score	split4_test_score	mean_test_score	std_test_score	rank_test_score
10	17.937267	0.343770	0.154968	0.026877	None ...	0.968582	0.970377	0.973617	0.003799	1		
7	17.345489	0.471931	0.200279	0.012649	90 ...	0.968582	0.973070	0.973258	0.003189	2		
11	25.724992	0.571347	0.151303	0.015432	None ...	0.970377	0.970377	0.973079	0.002280	3		
8	32.863902	0.466363	0.253013	0.034214	90 ...	0.966786	0.971275	0.972540	0.003801	4		
6	1.692553	0.161996	0.112161	0.012713	90 ...	0.964093	0.969479	0.970745	0.003757	5		

[5 rows x 15 columns]

	N-Gram:	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	...	split3_test_score	split4_test_score	mean_test_score	std_test_score	rank_test_score
11	94.849783	1.846754	0.388625	0.027227	None ...	0.944345	0.958628	0.949030	0.002512	1		
10	63.092635	1.635151	0.397301	0.033374	None ...	0.946140	0.947935	0.949030	0.001846	1		
9	7.595357	0.336770	0.496689	0.159212	None ...	0.934470	0.936266	0.938980	0.003221	3		
6	6.061857	1.331968	0.587655	0.204747	90 ...	0.925494	0.935368	0.928751	0.004737	4		
7	36.943602	0.880433	0.826666	0.213086	90 ...	0.921005	0.928187	0.928211	0.004306	5		

[5 rows x 15 columns]

	TF-IDF:	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	...	split3_test_score	split4_test_score	mean_test_score	std_test_score	rank_test_score
11	25.131840	0.457054	0.140176	0.018235	None ...	0.969479	0.973070	0.974514	0.003044	1		
7	15.952247	0.238358	0.196436	0.031546	90 ...	0.968582	0.971275	0.973976	0.003461	2		
8	31.410926	0.411874	0.261430	0.017647	90 ...	0.967684	0.971275	0.973258	0.003480	3		
10	17.234838	0.452019	0.147486	0.026116	None ...	0.968582	0.971275	0.973079	0.002959	4		
5	25.438713	0.638413	0.248299	0.013025	60 ...	0.966786	0.972172	0.971822	0.003098	5		

[5 rows x 15 columns]

Najlepszy mean_test_score:

Count: 0.97326

N-Gram: 0.94921

TF-IDF: 0.97523

Dyskusja wyników:

Na podstawie przeprowadzonych testów z użyciem GridSearchCV, najlepsze rezultaty klasyfikacji osiągnięto przy wykorzystaniu TF-IDF, gdzie mean_test_score był najwyższy. Niewiele gorzej wypadł CountVectorizer, co oznacza, że sama obecność słów w tekście również niesie dużą wartość informacyjną w kontekście klasyfikacji wiadomości jako spam lub ham. Najsłabszy wynik uzyskano dla reprezentacji N-Gram, co może wynikać z większego wymiaru cech i wprowadzenia rzadkich kombinacji wyrazów, które nie zawsze poprawiają jakość modelu. Podsumowując, TF-IDF okazał się najbardziej efektywnym podejściem, łącząc wysoką skuteczność z odpornością na szum w danych.